

Operating Systems

Processes

Yanyan Zhuang

Department of Computer Science

<http://www.cs.uccs.edu/~yzhuang>

Recap of the Last Class

- Computer hardware
 - Time-multiplexed
 - Space-multiplexed
- OS components
 - Process management
 - Memory management
 - File and storage management

Process

- Process == Program??

Process

- Definition
 - An instance of a program running on a computer
 - ▶ Process == Program in execution
 - An abstraction that supports running programs

Process


- Definition
 - An instance of a program running on a computer
 - ▶ Process == Program in execution
 - An abstraction that supports running programs
- Allows the OS to simplify
 - Resource allocation/management
 - OS manages resources and internal state of every process
 - ▶ Dynamic process state
 - Registers: PC, SP, ...
 - Memory: address space, code, stack, heap ...
 - I/O status: opened files ...

Program v.s Process

- Program != Process
 - Program = static code + data
 - Process = dynamic instantiation of code + data + files ...
- No 1:1 mapping
 - A program can invoke many processes
 - ▶ Running the same program twice
 - ▶ A program contains fork()

The Process Model


- Assume a single program counter
 - Each process in unique memory location
 - CPU switches back and forth from process to process



Multiprogramming of four programs

The Process Model

- Each process has their own flow of control (own logical program counter)
- Each time we switch processes, save the program counter of first process and restore the program counter of the second
 - Hiding the effects of interrupts and blocking system calls




Conceptual model of four independent, sequential processes



The Process Model


- Assume a single program counter
 - All processes make progress, but only one is active at any given time



(c)


The Process Model

One program counter




(a)

Four program counters



(b)



(c)

- CPU time can be allocated to different processes
 - ✓ OS normally offers **no timing or ordering guarantees**

Process Creation

- Principal events that cause process creation
 - System initialization
 - Execution of a process creation system call by an existing process
 - User request to create a new process
 - Initiation of a batch job

Process Creation

- Principal events that cause process creation
 - System initialization
 - Execution of a process creation system call by an existing process
 - User request to create a new process
 - Initiation of a batch job
- UNIX example
 - ***fork*** system call creates an **exact copy** of calling process
 - ▶ Same memory image, environment settings, and opened files
 - ▶ After fork, caller is parent, newly-created process is child

Process Creation

- Principal events that cause process creation
 - System initialization
 - Execution of a process creation system call by an existing process
 - User request to create a new process
 - Initiation of a batch job
- UNIX example
 - ***fork*** system call creates an **exact copy** of calling process
 - ▶ Same memory image, environment settings, and open files
 - ▶ After fork, caller is parent, newly-created process is child
 - Child process calls ***execve*** to change its memory image and run a new program

After fork()

- The child process returns executing at the exact same point after its parent called fork()
 - fork() returns **twice**: the new PID to the parent, and 0 to the child

```
pid= fork();
if (pid == 0) {
    /* I am the child (0: invalid PID) */
} else {
    /* I am the parent */
}
```

**Two processes execute the code!
(parent/child share same text)**

- All memory contents of parent/child are identical
- Both have the same files open at the same position (point to the same file objects)

Putting it Together

```
/* now create new process */
pid = fork();
char *const parmList[] = {"./Helloworld", NULL};
if (pid == 0) /* fork() returns 0 to the child process */
{
    sleep(1);
    printf("CHILD: My parent's PID: %d\n", getppid());
    execve("./Helloworld", parmList);
    printf("retval of Helloworld: %d\n", retval);
    exit(retval);
}
else /* fork() returns new pid to the parent process */
{
    printf("PARENT: my child PID: %d\n", pid);
    wait(&status);
    printf("PARENT: Child's exit code is: %d\n", WEXITSTATUS(status));
    exit(0);
}
```



Process Management System Calls

- fork: create a new process
 - Child is a **clone** of the parent
 - Shares **some** resources with the parent
 - exec: execute a new process image
 - Used in combination with fork
 - Has a family of calls, e.g., execve
 - exit: cause voluntary process termination
 - Exit status returned to the parent
 - wait: parent wait for child to finish
 - **wait(&status)**: a child's exit status will be stored in **status** upon return
 - kill: send a signal to a process (or group)
 - Can cause involuntary process termination
- Many of them have C lib call or command line wrappers




Process Termination


- Conditions which terminate processes
 - Normal exit (voluntary)
 - Error exit (voluntary)
 - Killed by another process (involuntary)

Process Hierarchies (Trees)

- Parent creates a child process, child processes can create its own process
- Forms a hierarchy
 - UNIX: a process and all its children and further descendants form a "process group"
 - *init*, a special process present in the boot image
 - Try: `pstree -h`



Process Life Cycle



To allocate CPU time, the OS needs to track **process states**:


- **Running**: process is currently executed by the CPU
- **Blocked**: process is waiting for some event
- **Ready**: process is waiting to be selected to run

Implementation of Processes

- Process table
 - One entry per process
 - Each entry is called a process control block (PCB)
- Process control block (PCB)
 - OS data structure containing data associated with processes
 - ▶ ID
 - ▶ Memory address space
 - ▶ Hardware registers (e.g., program counter)
 - ▶ Opened files
 - ▶

Linux Processes


- Process table: implemented as a linked list/
hashtable
 - Each element: a process descriptor of type **task_struct**
 - Dynamically allocated for
each process
 - Contains all info about a
specific process



Linux Processes


- Process descriptor (PCB)

- State
- Identifiers
- Scheduling info
- File system
- Virtual memory
- ...



Linux Process Descriptor


- State
 - **TASK_RUNNING**
 - ▶ Running
 - **TASK_INTERRUPTABLE**
 - ▶ Blocked
 - **EXIT_ZOMBIE**
 - ▶ Terminated by not deallocated
 - **EXIT_DEAD**
 - ▶ Completely terminated



Linux Process Descriptor (cont’)

- Identifiers
 - pid: PID of the process/thread
 -
- How to get the pointer to a specific process?
 - The **current** macro
 - The **init_task** macro
 - **find_task_by_vpid(pid_t pid)**

For projects,
don't use this



Use this

Kernel modules can only do this to get `task_struct`:

```
pid* pid_struct = find_get_pid(int pid);
task_struct* task = pid_task(pid_struct,PIDTYPE_PID);
```

Linux Process Descriptor (cont’)

- Files
 - `fs_struct`
 - ▶ file system information: root directory, current directory
 - `files_struct`
 - ▶ Information on opened files



Summary

- What is a process?
 - An instantiation of a program
- Program life cycle
 - Ready, running, blocked, new, terminated
- Process implementation
 - Process table, PCB
- Additional practice
 - Download Linux kernel source to your VM, find the following fields in structure `task_struct` (PCB) in `LINUX_SRC_FOLDER/include/linux/sched.h`
 - ▶ Program counter (try to google)
 - ▶ Stack pointer
 - ▶ Process ID
 - ▶ Opened file descriptors

