

Operating Systems

Threads

Yanyan Zhuang

Department of Computer Science



<http://www.cs.uccs.edu/~yzhuang>

Recap of the Last Class

- Processes
 - A program in execution
 - 5 (3)-state process model
 - Process control block
- Linux processes
 - The `task_struct` structure

Thread and Multithreading

- Process
 - Resource grouping and execution




So far, we assumed 1 thread of execution
(except fork)


Thread and Multithreading

- Process
 - Resource grouping and execution
- Thread (or multi-threaded execution)


- Each one is doing something
- They **share** the same data but look at different parts
- They have **private** state but can communicate easily
- They must coordinate!



An Illustration: from OS point of view



An Illustration: from OS point of view




Thread and Multithreading

- Process
 - Resource grouping and execution
- Thread
 - A program in execution without dedicated address space: *threads of the same process share address space*
 - Efficient communication
 - ▶ Inter-**thread** communication can be carried out via shared data objects within the shared address space
 - ▶ Inter-**process** communication usually requires other OS services
 - Efficient creation
 - ▶ Only create thread context

Processes v.s. Threads: A Closer Look

- Threads
 - No data segment or heap
 - Multiple can coexist in a process
 - Share code, data, heap, and I/O
 - Have own stack and registers
 - Inexpensive to create
 - Inexpensive context switching
 - Efficient communication
- Processes
 - Have data/code/heap
 - Include at least one thread
 - Have own address space, isolated from other processes
 - Expensive to create
 - Expensive context switching
 - IPC can be expensive

Thread Usage



A multithreaded Web server.

Thread interfaces in UNIX: POSIX threads

- UNIX systems
 - IEEE Portable Operating System Interface (POSIX)
 - Implementations of threads that adhere to this standard are referred to as POSIX threads, or Pthreads

Pthread function

- `int pthread_create(pthread_t *thread,
const pthread_attr_t *attr,
void *(*start_routine) (void *),
void *arg)`
 - Create a new thread, with attributes attr (attributes can include scheduling policies, stack size, etc.)
 - The thread is created by executing `start_routine` with `arg` as the only argument
 - Upon success, stores the ID of the thread in the location referenced by `thread`

An Example


```
void *my_thread(void *arg)
{
    int *tid = (int *)arg;
    printf("Hello from child thread: %d\n", *tid);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t threads[NR_THREADS];
    for (i = 1; i < NR_THREADS; i++) {
        printf("In main: creating thread %ld\n", i);
        tid[i] = i;
        pthread_create(&threads[i], &a, my_thread, &tid[i]);
    }
}
```



Implementing Threads in User-Space

- User-level threads: the kernel knows nothing about them



A user-level threads package

- OS thinks there's only a single-threaded process
- Threads in same process don't involve multiplexing between processes so no kernel privilege required

User-level Thread - Discussions

- Advantages
 - No OS thread-support needed
 - Lightweight: thread switching vs. process switching
 - Local procedure (no mode switch) vs. system call (trap to kernel)
 - Each has its own customized scheduling algorithms
 - `thread_yield()`

User-level Thread - Discussions


- Advantages
 - No OS thread-support needed
 - Lightweight: thread switching vs. process switching
 - Local procedure (no mode switch) vs. system call (trap to kernel)
 - Each has its own customized scheduling algorithms
 - `thread_yield()`
- Disadvantages
 - How blocking system calls implemented? Called by a thread?
 - Goal: to allow each thread to use blocking calls, but to prevent one blocked thread from affecting the others
 - How to deal with page faults?
 - How to stop a thread from running forever?
 - No clock interrupts in a single process

Implementing Threads in the Kernel

- Kernel-level threads: when a thread blocks, kernel reschedules another thread

- Threads known to OS
 - Scheduled by the scheduler
- Slow
 - Trap into the kernel mode
- Expensive to create and switch
 - Less expensive if in the same process
 - Registers, PC, stack pointer need to be created/changed
 - Not the memory info

Any problems?




What happens when forking a multithreaded process



Hybrid Implementations

- Use kernel-level threads and then *multiplex* user-level threads onto some or all of the kernel-level threads
- Multiplexing user-level threads onto kernel-level threads
- Enjoy the benefits of user and kernel level threads

Too complex!



Multiplexing user-level threads onto kernel-level threads

Threading Models

- N:1 (User-level threading)
 - N user threads that look like 1 to the OS kernel
- 1:1 (Kernel-level threading)
 - Each user-thread is paired with a kernel-thread (aka native thread)
- M:N (Hybrid threading)
 - Solaris

Threads in Linux

- Thread control block (TCB)
 - The `thread_struct` structure
 - Includes registers and processor-specific context
- Linux treats threads like processes
 - Use `clone()` to create threads instead of using `fork()`
 - `clone()` is usually not called directly, but from some threading libraries

Summary

- Processes v.s. threads?
- Why threads?
 - Concurrency + lightweight
- Threading models
 - N:1, 1:1, M:N