

Operating Systems

Synchronization

Yanyan Zhuang

Department of Computer Science

<http://www.cs.uccs.edu/~yzhuang>

Inter-Process Communication (IPC)

- Fundamental issues:
 - How to make sure two or more processes do not get in each other's way when engaging in critical activities
 - How to maintain proper sequencing when dependencies present

Inter-Process Communication (IPC)

- Fundamental issues:
 - How to make sure two or more processes do not get in each other's way when engaging in critical activities
 - How to maintain proper sequencing when dependencies present

Inter-thread communication?


Same problems exist & same solutions apply

Race Conditions

- Race conditions: when two or more processes/threads are reading or writing some *shared data* and the final results depend on who runs precisely when
 - Interrupts, interleaved operations/execution


```
free_slot = in;  
Dir[free_slot] = X;  
in ++;
```

```
free_slot = in;  
Dir[free_slot] = Y;  
in ++;
```



Race Conditions

- Race conditions: when two or more processes/threads are reading or writing some *shared data* and the final results depend on who runs precisely when
 - Interrupts, interleaved operations/execution




Mutual Exclusion and Critical Regions

- Mutual exclusion: makes sure if one process is using a shared variable or file, the other processes will be excluded from doing the same thing
 - Main challenge/issue to OS: to design appropriate primitive operations for achieving mutual exclusion in user space
- Critical regions: the part of the program where the shared resource is accessed

Mutual Exclusion and Critical Regions

- Mutual exclusion: makes sure if one process is using a shared variable or file, the other processes will be excluded from doing the same thing
 - Main challenge/issue to OS: to design appropriate primitive operations for achieving mutual exclusion in user space
- Critical regions: the part of the program where the shared resource is accessed
- Four conditions to provide mutual exclusion
 - No two processes simultaneously in critical region
 - No assumptions made about speeds or numbers of CPUs
 - No process running outside its critical region may block another process
 - No process must wait forever to enter its critical region

Mutual Exclusion Using Critical Regions



Mutual exclusion using critical regions

Mutual Exclusion with Busy Waiting

- Disabling interrupts:
 - OS technique, not users'
 - Multi-CPU?
- Lock variables:
 - Test-set is a two-step process, not atomic
- Busy waiting:
 - Continuously testing a variable until some value appears (*spin lock*)

Busy Waiting: Strict Alternation

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Proposed *strict alternation* solution to critical region problem

(a) Process 0.

(b) Process 1.

Busy Waiting: Strict Alternation

```
while (TRUE) {  
    while (turn != 0);      /* loop */;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1);      /* loop */;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Proposed *strict alternation* solution to critical region problem

(a) Process 0.

(b) Process 1.



Busy Waiting: Strict Alternation

```
while (TRUE) {  
    while (turn != 0);      /* loop */;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1);      /* loop */;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Proposed *strict alternation* solution to critical region problem

(a) Process 0.

(b) Process 1.

What if P1's noncritical_region() has lots more work than P0's?

Taking turns is not a good idea when one is much slower than the other

Busy Waiting: TSL

- TSL (Test and Set Lock)
 - Indivisible (**atomic**) operation, how? Hardware (multi-processor)
 - How to use TSL to prevent two processes from simultaneously entering their critical regions? 0 can enter; 1 can't enter

enter_region:

TSL REGISTER,LOCK	copy lock to register and set lock to 1
CMP REGISTER,#0	was lock zero?
JNE enter_region	if it was non zero, lock was set, so loop
RET return to caller; critical region entered	

leave_region:

MOVE LOCK,#0	store a 0 in lock
RET return to caller	

Entering and leaving a critical region using the TSL instruction



Busy Waiting: Pros and Cons

- Cons
 - Wastes CPU cycles
 - For single-core system, user-level threads
 - ▶ T1 waiting for T2 to change lock to 0, but T2 never gets to run when T1 is running
 - ▶ Priority inversion
 - 2 processes on 1 CPU. Process H with high priority and L with low priority, L is in its critical region and H becomes ready; *does L have chance to leave critical region?*
- Pros
 - Avoids expensive context switch when critical region is *very short*
 - ▶ Sleep/wakeup (an alternative to busy waiting) requires context switch
 - Time to put a thread to sleep/wake up might exceed the time a thread has actually slept



Sleep and Wakeup

- Issue I: avoid CPU-costly busy waiting
- Issue II: avoid **priority inversion problem**
- Some IPC primitives that **block** instead of wasting CPU time when they are not allowed to enter their critical regions
 - Sleep and wakeup

Sleep and Wakeup – Producer-Consumer Problem

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                    /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        item = produce_item();                     /* generate next item */
        if (count == N) sleep();                   /* if buffer is full, go to sleep */
        insert_item(item);                        /* put item in buffer */
        count = count + 1;                         /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);          /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        if (count == 0) sleep();                   /* if buffer is empty, got to sleep */
        item = remove_item();                     /* take item out of buffer */
        count = count - 1;                        /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);      /* was buffer full? */
        consume_item(item);                      /* print item */
    }
}
```



Sleep and Wakeup – Producer-Consumer Problem

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                    /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        item = produce_item();                     /* generate next item */
        if (count == N) sleep();                   /* if buffer is full, go to sleep */
        insert_item(item);                        /* put item in buffer */
        count = count + 1;                         /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);          /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        if (count == 0) sleep();                   /* if buffer is empty, got to sleep */
        item = remove_item();                     /* take item out of buffer */
        count = count - 1;                        /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);      /* was buffer full? */
        consume_item(item);                      /* print item */
    }
}
```

Producer: buffer full → sleep, wakeup when item removed
Consumer: buffer empty → sleep, wakeup when item inserted



Sleep and Wakeup – Producer-Consumer Problem

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                    /* number of items in the buffer */

void producer(void)
{
    int item;      Q1: What if the wakeup signal sent to a non-sleep process?

    while (TRUE) {
        item = produce_item();                  /* repeat forever */
        if (count == N) sleep();                /* generate next item */
        insert_item(item);                    /* if buffer is full, go to sleep */
        count = count + 1;                   /* put item in buffer */
        if (count == 1) wakeup(consumer);     /* increment count of items in buffer */
                                            /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;      Interrupt!

    while (TRUE) {                          /* repeat forever */
        if (count == 0) sleep();            /* if buffer is empty, got to sleep */
        item = remove_item();             /* take item out of buffer */
        count = count - 1;                /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);              /* print item */
    }
}
```



Sleep and Wakeup – Producer-Consumer Problem

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                    /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();                    /* repeat forever */
        if (count == N) sleep();                  /* generate next item */
        insert_item(item);                      /* if buffer is full, go to sleep */
        count = count + 1;                      /* put item in buffer */
        if (count == 1) wakeup(consumer);        /* increment count of items in buffer */
                                                /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep();                /* repeat forever */
        item = remove_item();                  /* if buffer is empty, got to sleep */
        count = count - 1;                     /* take item out of buffer */
                                                /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);   /* was buffer full? */
        consume_item(item);                  /* print item */
    }
}
```

Q2: wakeup waiting bit
wakeup() sent to a process that is still awake,
wakeup waiting bit is set

Semaphores and P&V Operations

- Semaphores: a variable to indicate the # of pending wakeups (Dijkstra)
- *Down operation (request): lock/sleep*
 - Checks if a semaphore is > 0,
 - if so, it decrements the value and just continue
 - Otherwise (==0), the process is put to sleep



Semaphores and P&V Operations

- Semaphores: a variable to indicate the # of pending wakeups (Dijkstra)
- *Down* operation (request): **lock/sleep**
 - Checks if a semaphore is > 0,
 - if so, it decrements the value and just continue
 - Otherwise (==0), the process is put to sleep
- *Up* operation (release): **unlock/wakeup**
 - Increments the value of the semaphore
 - But if one or more processes are sleeping on the semaphore (==0), one of them is chosen (randomly) and wake up (**semaphore will still be 0**)



Semaphores and P&V Operations

- Semaphores: a variable to indicate the # of pending wakeups (Dijkstra)
- *Down* operation (request): **lock/sleep**
 - Checks if a semaphore is > 0,
 - if so, it decrements the value and just continue
 - Otherwise (==0), the process is put to sleep
- *Up* operation (release): **unlock/wakeup**
 - Increments the value of the semaphore
 - But if one or more processes are sleeping on the semaphore (==0), one of them is chosen (randomly) and wake up (**semaphore will still be 0**)
- P & V operations are **atomic**, how to implement?
 - Single CPU: system calls, OS disabling interrupts temporarily
 - Multiple CPUs: TSL



The Producer-consumer Problem w/ Semaphores

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */

For mutual exclusion
and synchronization

Binary semaphores: if each process does a down before entering its critical region and an up just leaving it, mutual exclusion is achieved



The Producer-consumer Problem w/ Semaphores

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

P: down(&empty) →

empty = N → empty = N-1

C: down(&full) → full=0
→ sleep on full

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */
```

For mutual exclusion
and synchronization

Binary semaphores: if each process does a down before entering its critical region and an up just leaving it, mutual exclusion is achieved



The Producer-consumer Problem w/ Semaphores

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
```

```
{    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}
```

```
void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

P: down(&empty) →

empty = N → empty = N-1

C: down(&full) → full=0
→ sleep on full

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
```

```
/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */
```

Binary semaphores: if each process does a down before entering its critical region and an up just leaving it, mutual exclusion is achieved

For mutual exclusion
and synchronization

P: up(&full) → C sleeps
on full → wakeup C, full=0

```
/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */
```

C: up(&empty) → empty =
N-1 → empty = N



Mutexes

- **Mutex:**

- A variable that can be in one of two states: unlocked or locked
- A simplified version of the semaphores [0, 1]

mutex_lock:

```
TSL REGISTER,MUTEX  
CMP REGISTER,#0  
JZE ok  
CALL thread_yield  
JMP mutex_lock
```

ok: RET | return to caller; critical region entered

| copy mutex to register and set mutex to 1
| was mutex zero?
| if it was zero, mutex was unlocked, so return
| mutex is busy; schedule another thread
| try again later

mutex_unlock:

```
MOVE MUXTEX,#0  
RET | return to caller
```

| store a 0 in mutex

Give others chance to run;
What about mutex_trylock()? Either gets the lock
or returns a code for failure → does not block



Mutexes – User-space Multi-threading

- What is a key difference between *mutex_lock* and *enter_region* in multi-threading and multi-processing?
 - For single-core, user-space multi-threading, a thread has to allow other thread to run and release the lock so as to enter its critical region, which is impossible with busy waiting *enter_region*

enter_region:

TSL REGISTER,LOCK	copy lock to register and set lock to 1
CMP REGISTER,#0	was lock zero?
JNE enter_region	if it was non zero, lock was set, so loop
RET return to caller; critical region entered	

leave_region:

MOVE LOCK,#0	store a 0 in lock
RET return to caller	

Two *processes* entering and leaving a critical region using the TSL instruction



When to Use Busy-waiting/Sleep-wakeup

- On a single-core system
 - Using busy-waiting makes no sense
 - ▶ No other thread can run, locks won't be unlocked
 - Use sleep-wakeup instead
- On a multi-core systems, (only when) locks are held for a very short time
 - Time for putting threads to sleep/waking up might decrease runtime performance
 - Use busy-waiting instead

Semaphores/Mutexes easy?

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = 0;
semaphore full = N
```

```
void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}
```

```
void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
```

Swap the order of mutex and empty
If the buffer is **full**

```
/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */
```

Must be careful using semaphores!

```
/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */
```



Monitors (1)

- Monitor: a higher-level synchronization primitive
 - Only **one process can be active in a monitor** at any instant, *with compiler's help*; thus, how about putting all the critical regions into monitor procedures for *mutual exclusion*?

```
monitor example
    integer i;
    condition c;

    procedure producer();
        .
        .
        .
    end;

    procedure consumer();
        .
        .
        .
    end;
end monitor;
```

Processes can call procedures in a monitor, but not directly access monitor's internal data

Programming-language construct: compiler knows they are special and handle calls differently

Monitors (1)

- Monitor: a higher-level synchronization primitive
 - Only **one process can be active in a monitor** at any instant, *with compiler's help*; thus, how about putting all the critical regions into monitor procedures for *mutual exclusion*?

```
monitor example
    integer i;
    condition c;

    procedure producer();
        .
        .
        .
    end;

    procedure consumer();
        .
        .
        .
    end;
end monitor;
```

But, how processes block when they cannot proceed?

Condition variables, and two operations: *wait()* and *signal()*

Monitors (2)

Wakeup and sleep signals can be lost, but not wait and signal, why?

```
monitor ProducerConsumer
    condition full, empty;
    integer count;
    procedure insert(item: integer);
begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
end;
function remove: integer;
begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
end;
count := 0;
end monitor;
```

```
procedure producer;
begin
    while true do
begin
    item = produce_item;
    ProducerConsumer.insert(item)
end
end;
procedure consumer;
begin
    while true do
begin
    item = ProducerConsumer.remove;
    consume_item(item)
end
end;
```



Monitors (2)

Wakeup and sleep signals can be lost, but not wait and signal, why?

```
monitor ProducerConsumer
```

```
    condition full, empty;
```

```
    integer count;
```

```
    procedure insert(item: integer);
```

```
begin Consumer not allowed in
```

```
    if count = N then wait(full);
```

```
    insert_item(item);
```

```
    count := count + 1;
```

```
    if count = 1 then signal(empty)
```

```
end;
```

```
function remove: integer;
```

```
begin Producer not allowed in
```

```
    if count = 0 then wait(empty);
```

```
    remove = remove_item;
```

```
    count := count - 1;
```

```
    if count = N - 1 then signal(full)
```

```
end;
```

```
    count := 0;
```

```
end monitor;
```

```
procedure producer;
```

```
begin
```

```
    while true do
```

```
        begin
```

```
            item = produce_item;
```

```
            ProducerConsumer.insert(item)
```

```
        end
```

```
    end;
```

```
procedure consumer;
```

```
begin
```

```
    while true do
```

```
        begin
```

```
            item = ProducerConsumer.remove;
```

```
            consume_item(item)
```

```
        end
```

```
    end;
```

Mutual exclusion guarantees e.g., producer will be able to complete wait() without switching to consumer before wait() completes



Monitors (3)

- Pros
 - Make mutual exclusion automatic
 - Make parallel programming less error-prone
- Cons
 - Compiler support
 - ▶ Keyword **synchronized** in Java



All previous methods...

- Work on a single computer
 - Assuming data can be shared in some way
- Data exchange between machines?
 - Message passing
 - ▶ Inter-process communication using two primitives, send/receive
 - ▶ `send(destination, &message);`
 - ▶ `receive(source, &message);`
 - ▶ System calls rather than language constructs

Message Passing

```
#define N 100                                     /* number of slots in the buffer */

void producer(void)
{
    int item;
    message m;                                    /* message buffer */

    while (TRUE) {
        item = produce_item( );
        receive(consumer, &m);
        build_message(&m, item);
        send(consumer, &m);
    }
}


void consumer(void)
{
    int item, i;
    message m;

    for (i = 0; i < N; i++) send(producer, &m); /* send N empties */
    while (TRUE) {
        receive(producer, &m);                  /* get message containing item */
        item = extract_item(&m);                /* extract item from message */
        send(producer, &m);                    /* send back empty reply */
        consume_item(item);                  /* do something with the item */
    }
}
```

The producer-consumer problem with N messages

Class IPC Problems: Dining Philosophers


- Philosophers eat/think
- Eating needs 2 forks
- Pick one fork at a time
- How to prevent deadlock & starvation
 - *Deadlock*: processes are blocked on some resource
 - *Starvation*: waiting for scheduler, no progress can be made



The problem is useful for modeling processes that are competing for exclusive access to a limited number of resources, such as I/O devices

Class IPC Problems: Dining Philosophers

- Philosophers eat/think
- Eating needs 2 forks
- Pick one fork at a time
- How to prevent deadlock & starvation
 - *Deadlock*: processes are blocked on some resource
 - *Starvation*: waiting for scheduler, no progress can be made



The problem is useful for modeling processes that are competing for exclusive access to a limited number of resources, such as I/O devices

Dining Philosophers (2)

```
#define N 5                                     /* number of philosophers */

void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */

{
    while (TRUE) {
        think();                                /* philosopher is thinking */
        take_fork(i);                            /* take left fork */
        take_fork((i+1) % N);                   /* take right fork; % is modulo operator */
        eat();                                   /* yum-yum, spaghetti */
        put_fork(i);                            /* put left fork back on the table */
        put_fork((i+1) % N);                   /* put right fork back on the table */
    }
}
```

Dining Philosophers (2)

```
#define N 5                                     /* number of philosophers */  
  
void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */  
{  
    while (TRUE) {  
        think();                                /* philosopher is thinking */  
        take_fork(i);                            /* take left fork */  
        take_fork((i+1) % N);                   /* take right fork; % is modulo operator */  
        eat();                                   /* yum-yum, spaghetti */  
        put_fork(i);                            /* put left fork back on the table */  
        put_fork((i+1) % N);                   /* put right fork back on the table */  
    }  
}
```

A **non-solution** to the dining philosophers problem

What happens if all philosophers pick up their left forks simultaneously?

Dining Philosophers (2)

```
#define N 5                                     /* number of philosophers */  
  
void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */  
{  
    while (TRUE) {  
        think();                                /* philosopher is thinking */  
        take_fork(i);                            /* take left fork */  
        take_fork((i+1) % N);                   /* take right fork; % is modulo operator */  
        eat();                                   /* yum-yum, spaghetti */  
        put_fork(i);                            /* put left fork back on the table */  
        put_fork((i+1) % N);                   /* put right fork back on the table */  
    }  
}  
} A non-solution to the dining philosophers problem
```

What happens if all philosophers pick up their left forks simultaneously?

Dining Philosophers (2)

```
#define N 5                                     /* number of philosophers */  
  
void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */  
{  
    while (TRUE) {  
        think();                                /* philosopher is thinking */  
        take_fork(i);                            /* take left fork */  
        take_fork((i+1) % N);                   /* take right fork; % is modulo operator */  
        eat();                                   /* yum-yum, spaghetti */  
        put_fork(i);                            /* put left fork back on the table */  
        put_fork((i+1) % N);                   /* put right fork back on the table */  
    }  
}  
}                                              A non-solution to the dining philosophers problem
```

What if *down* and *up* on *mutex* before acquiring/replacing a fork?
Only one can be eating at any instant...



Dining Philosophers (3): Solution part1

```
#define N          5           /* number of philosophers */
#define LEFT        (i+N-1)%N   /* number of i's left neighbor */
#define RIGHT       (i+1)%N    /* number of i's right neighbor */
#define THINKING    0           /* philosopher is thinking */
#define HUNGRY      1           /* philosopher is trying to get forks */
#define EATING      2           /* philosopher is eating */
typedef int semaphore;
int state[N];
semaphore mutex = 1;           /* mutual exclusion for critical regions */
semaphore s[N];               /* one semaphore per philosopher */

void philosopher(int i)         /* i: philosopher number, from 0 to N-1 */
{
    while (TRUE) {              /* repeat forever */
        think();                /* philosopher is thinking */
        take_forks(i);          /* acquire two forks or block */
        eat();                   /* yum-yum, spaghetti */
        put_forks(i);            /* put both forks back on table */
    }
}
```



Dining Philosophers (4): Solution part2

```
void take_forks(int i)                                /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);
    state[i] = HUNGRY;
    test(i);
    up(&mutex);
    down(&s[i]);
}

void put_forks(i)                                     /* i: philosopher number, from 0 to N-1 */
{
    down(&mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    up(&mutex);
}

void test(i)                                         /* i: philosopher number, from 0 to N-1 */
{
    if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
        state[i] = EATING;
        up(&s[i]);
    }
}
```



Readers and Writers Problem

- Models access to a database
 - Acceptable to have multiple processes reading the database at same time
 - One process writing to the database: no other processes may access (even readers)

Readers and Writers Problem

```
typedef int semaphore;
semaphore mutex = 1;
semaphore db = 1;
int rc = 0;

/* use your imagination */
/* controls access to 'rc' */
/* controls access to the database */
/* # of processes reading or wanting to */

void reader(void)
{
    while (TRUE) {
        down(&mutex);
        rc = rc + 1;
        if (rc == 1) down(&db);
        up(&mutex);
        read_data_base();
        down(&mutex);
        rc = rc - 1;
        if (rc == 0) up(&db);
        up(&mutex);
        use_data_read();
    }
}

void writer(void)
{
    while (TRUE) {
        think_up_data();
        down(&db);
        write_data_base();
        up(&db);
    }
}
```



Readers and Writers Problem

```
typedef int semaphore;
semaphore mutex = 1;
semaphore db = 1;
int rc = 0;

/* use your imagination */
/* controls access to 'rc' */
/* controls access to the database */
/* # of processes reading or wanting to */

void reader(void)
{
    while (TRUE) {
        /* repeat forever */
        down(&mutex);
        rc = rc + 1;
        if (rc == 1) down(&db);
        up(&mutex);
        read_data_base();
        down(&mutex);
        rc = rc - 1;
        if (rc == 0) up(&db);
        up(&mutex);
        use_data_read();
    }
}

void writer(void)
{
    while (TRUE) {
        /* repeat forever */
        think_up_data();
        down(&db);
        write_data_base();
        up(&db);
    }
}
```



Readers and Writers Problem

```
typedef int semaphore;
semaphore mutex = 1;
semaphore db = 1;
int rc = 0;

/* use your imagination */
/* controls access to 'rc' */
/* controls access to the database */
/* # of processes reading or wanting to */

void reader(void)
{
    while (TRUE) {
        /* repeat forever */
        down(&mutex);
        /* get exclusive access to 'rc' */
        rc = rc + 1;
        if (rc == 1) down(&db);
        up(&mutex);
        read_data_base();
        down(&mutex);
        rc = rc - 1;
        if (rc == 0) up(&db);
        up(&mutex);
        use_data_read();
    }
}

void writer(void)
{
    while (TRUE) {
        /* repeat forever */
        think_up_data();
        /* noncritical region */
        down(&db);
        write_data_base();
        up(&db);
    }
}
```



Readers and Writers Problem

- Problem
 - Additional readers admitted as they come along
 - As long as there's a steady supply of readers, writer will be suspended
- Solution
 - When a reader arrives, if a writer is waiting, the reader is suspended behind the writer

Summary

- Race condition, mutual exclusion
- Classic IPC problems
- Additional practice
 - Read Linux documentation:
`LINUX_SRC/Documentation/spinlocks.txt`
 - Find the implementation of `down` and `up` in
`LINUX_SRC/kernel/semaphore.c`
 - Spinlock v.s. Mutex:
<http://stackoverflow.com/questions/5869825/when-should-one-use-a-spinlock-instead-of-mutex>

